

Chapter 25 Binary Search Trees






Objectives

- To design and implement a binary search tree (§25.2).
- To represent binary trees using linked data structures (§25.2.1).
- To search an element in binary search tree (§25.2.2).
- To insert an element into a binary search tree (§25.2.3).
- To traverse elements in a binary tree (§25.2.4).
- To delete elements from a binary search tree (§25.3).
- To display binary tree graphically (§25.4).
- To create iterators for traversing a binary tree (§25.5).
- To implement Huffman coding for compressing data using a binary tree (§25.6).



Binary Trees

A list, stack, or queue is a linear structure that consists of a sequence of elements. A binary tree is a hierarchical structure. It is either empty or consists of an element, called the *root*, and two distinct binary trees, called the *left subtree* and *right subtree*.



See How a Binary Search Tree Works

<http://www.cs.armstrong.edu/liang/animation/web/BST.html>



Binary Search Tree Animation

www.cs.armstrong.edu/liang/animation/web/BST.html

BST Animation by Y. Daniel Liang

Enter an integer key and click the Search button to search the key in the tree. Click the Insert button to insert the key into the tree. Click the Remove button to remove the key from the tree. For the best display, use integers between 0 and 99.

```
graph TD; 5((5)) --- 4((4)); 5 --- 15((15)); 15 --- 52((52));
```

Enter a key: 52 Search Insert Remove Show Inorder Show Preorder Show Postorder


Binary Tree Terms

The root of left (right) subtree of a node is called a *left (right) child* of the node. A node without children is called a *leaf*. A special type of binary tree called a *binary search tree* is often useful. A binary search tree (with no duplicate elements) has the property that for every node in the tree the value of any node in its left subtree is less than the value of the node and the value of any node in its right subtree is greater than the value of the node. The binary trees in Figure 25.1 are all binary search trees. This section is concerned with binary search trees.




Representing Binary Trees

A binary tree can be represented using a set of linked nodes. Each node contains a value and two links named *left* and *right* that reference the left child and right child, respectively, as shown in Figure 25.2.



```
class TreeNode<E> {  
    E element;  
    TreeNode<E> left;  
    TreeNode<E> right;  
  
    public TreeNode(E o) {  
        element = o;  
    }  
}
```



Searching an Element in a Binary Search Tree

```
public boolean search(E element) {  
    TreeNode<E> current = root; // Start from the root  
    while (current != null)  
        if (element < current.element) {  
            current = current.left; // Go left  
        }  
        else if (element > current.element) {  
            current = current.right; // Go right  
        }  
        else // Element matches current.element  
            return true; // Element is found  
    return false; // Element is not in the tree  
}
```



Inserting an Element to a Binary Search Tree

If a binary tree is empty, create a root node with the new element. Otherwise, locate the parent node for the new element node. If the new element is less than the parent element, the node for the new element becomes the left child of the parent. If the new element is greater than the parent element, the node for the new element becomes the right child of the parent. Here is the algorithm:




Inserting an Element to a Binary Tree

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.




Trace Inserting 101 into the following tree

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



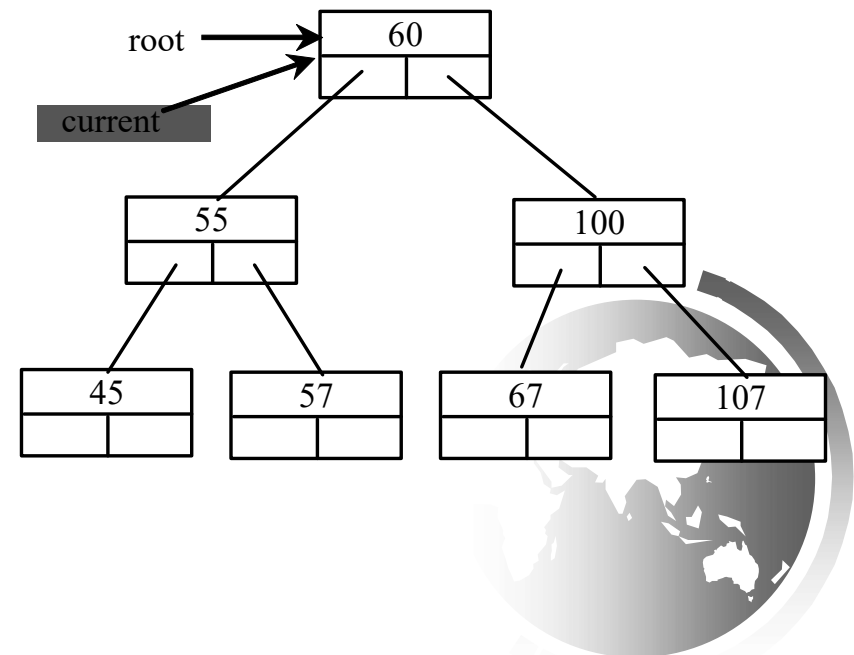
Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.




Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.




Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.




Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



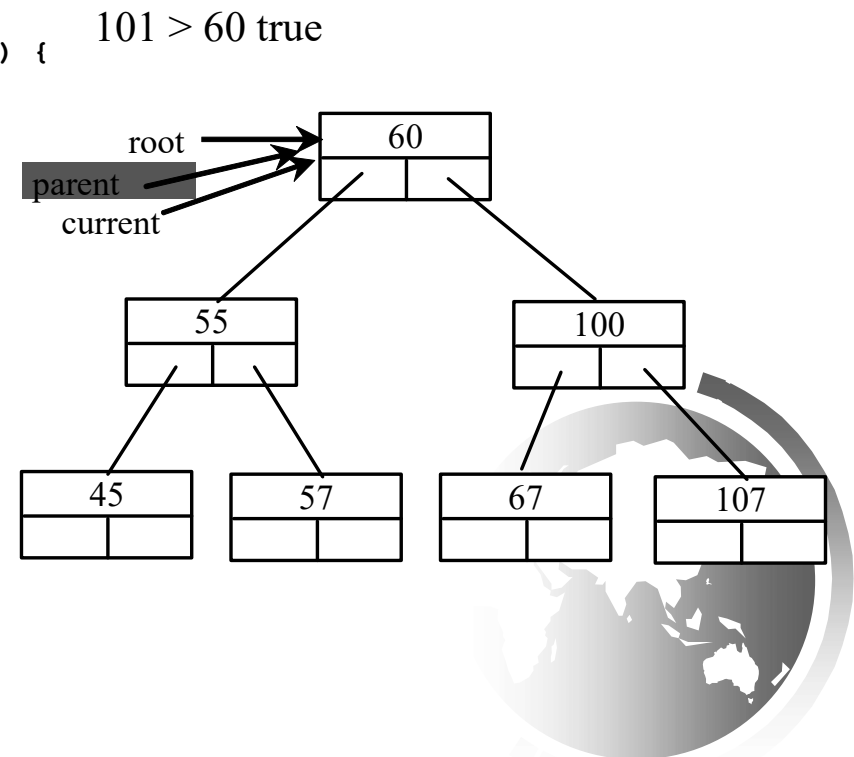
Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



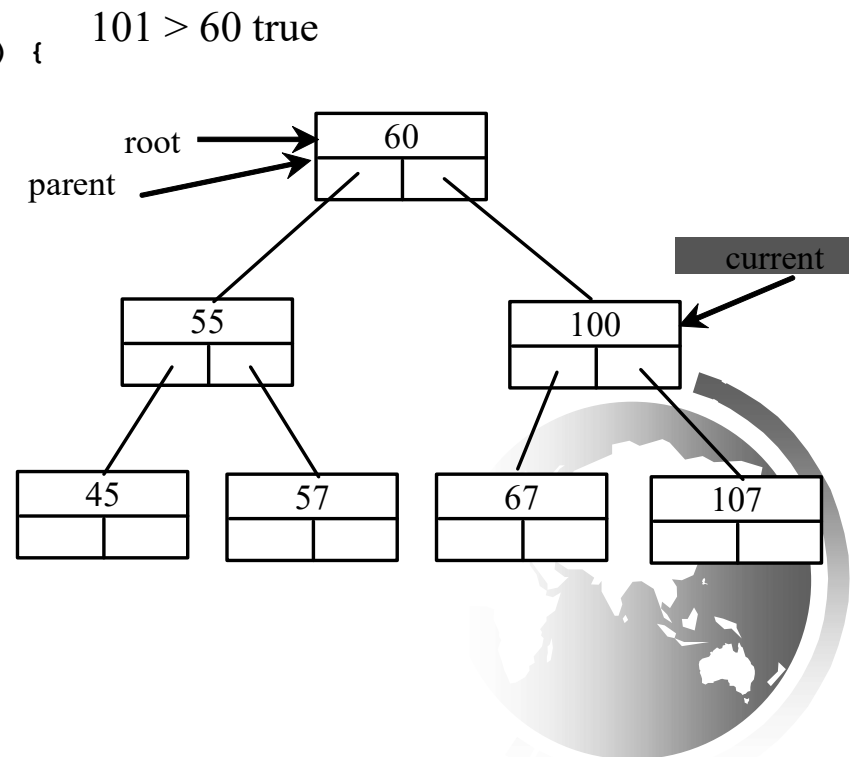
Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



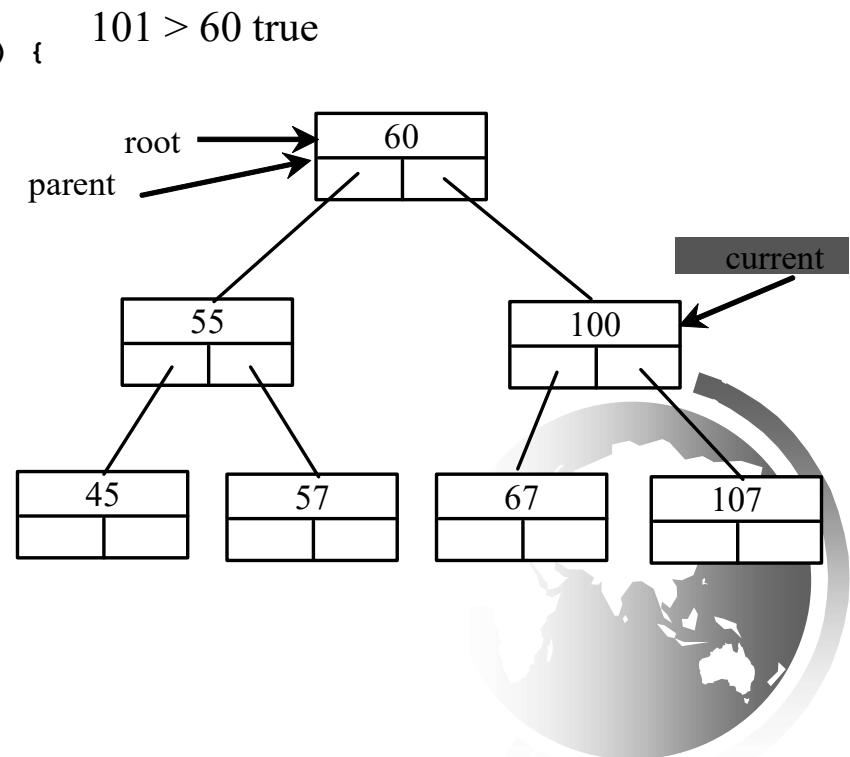
Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



Trace Inserting 101 into the following tree, cont.

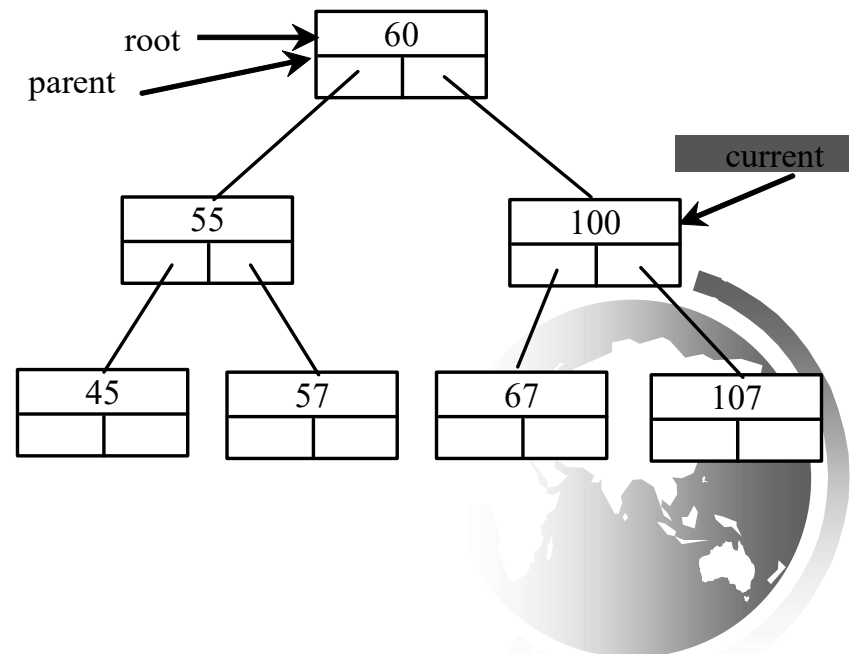
```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.

101 > 100 false



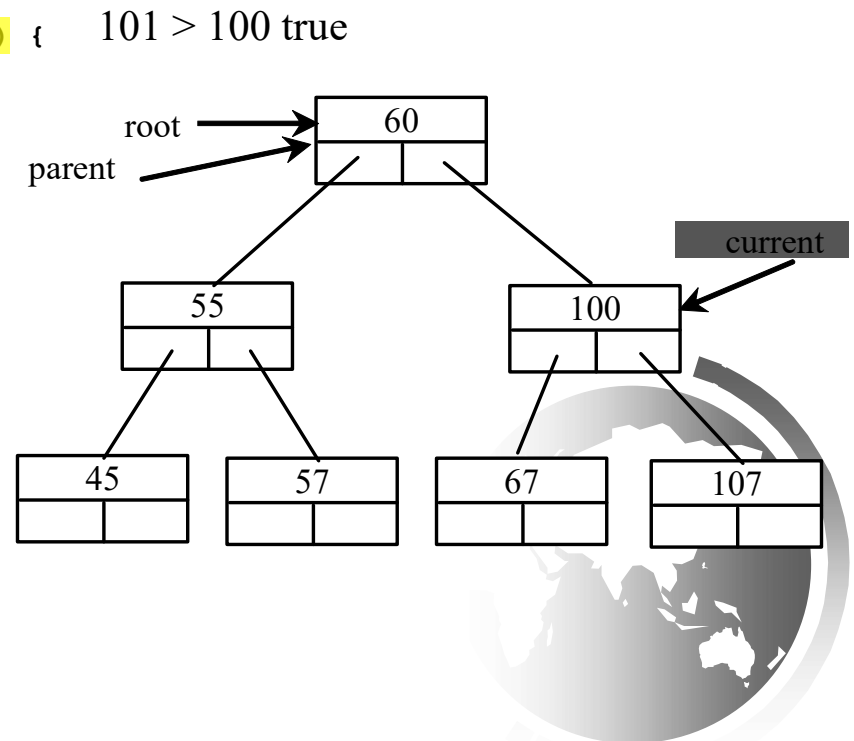
Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.




Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



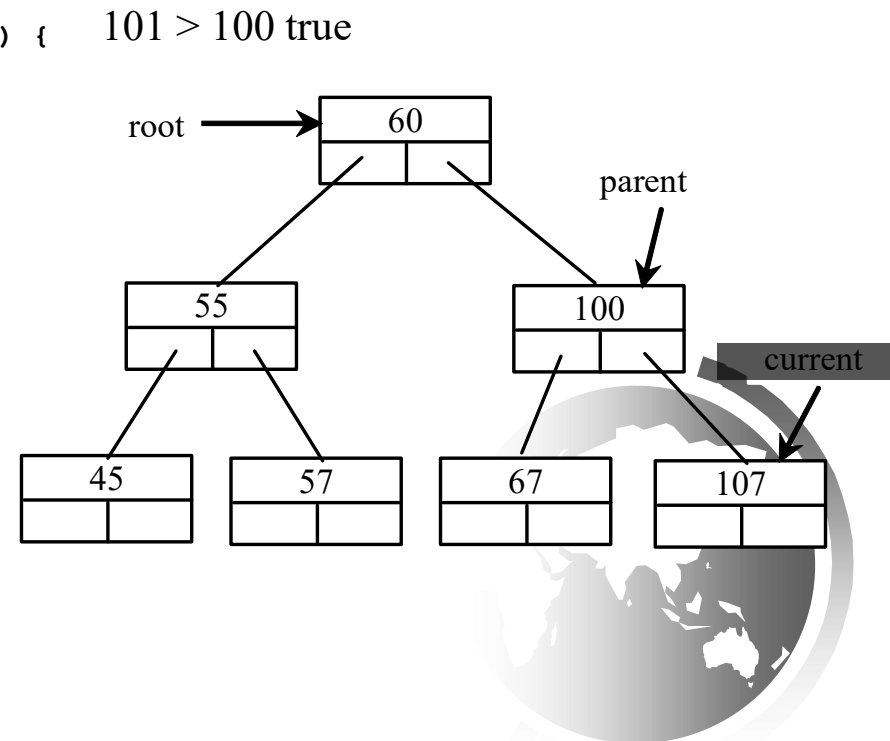
Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



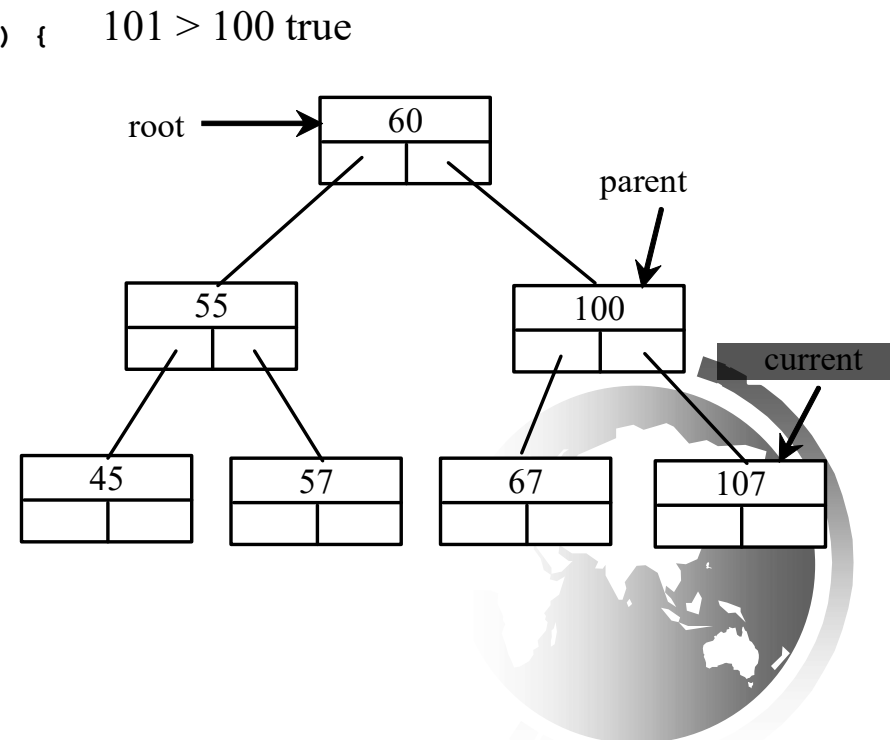
Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



Trace Inserting 101 into the following tree, cont.


```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.

101 < 107 true



Trace Inserting 101 into the following tree, cont.

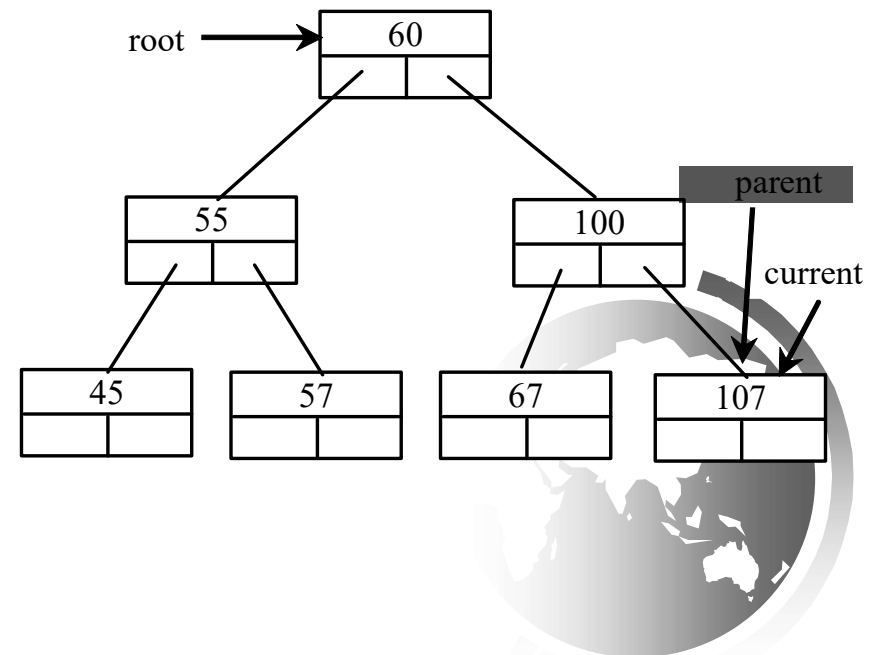
```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.

101 < 107 true



Trace Inserting 101 into the following tree, cont.


```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.

101 < 107 true



Since current.left is null, current becomes null


Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)           current is null now
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
        else
            return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.



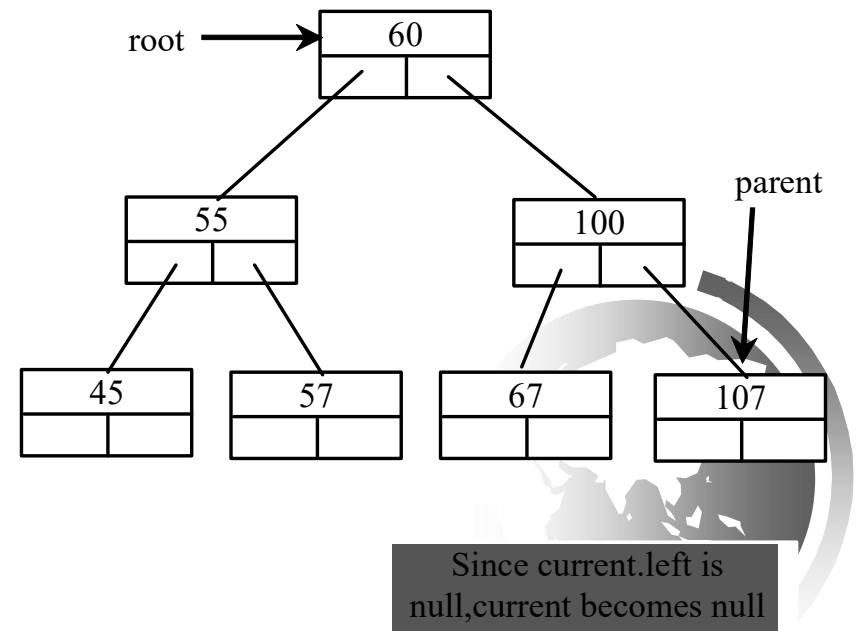
Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.




Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.




Trace Inserting 101 into the following tree, cont.

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```

Insert 101 into the following tree.

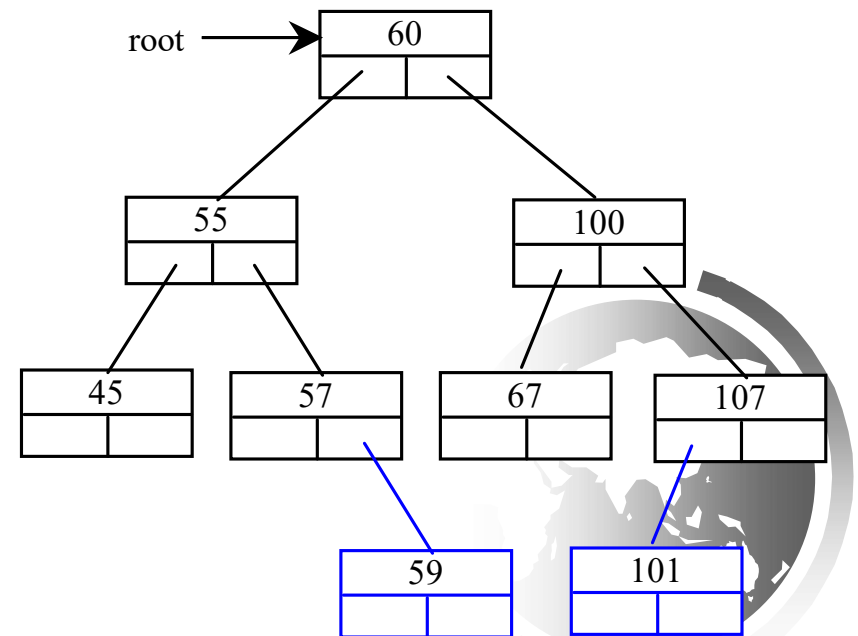


Inserting 59 into the Tree

```
if (root == null)
    root = new TreeNode(element);
else {
    // Locate the parent node
    current = root;
    while (current != null)
        if (element value < the value in current.element) {
            parent = current;
            current = current.left;
        }
        else if (element value > the value in current.element) {
            parent = current;
            current = current.right;
        }
    }
    else
        return false; // Duplicate node not inserted

    // Create the new node and attach it to the parent node
    if (element < parent.element)
        parent.left = new TreeNode(element);
    else
        parent.right = new TreeNode(element);

    return true; // Element inserted
}
```



Tree Traversal

Tree traversal is the process of visiting each node in the tree exactly once. There are several ways to traverse a tree. This section presents *inorder*, *preorder*, *postorder*, *depth-first*, and *breadth-first* traversals.

The inorder traversal is to visit the left subtree of the current node first recursively, then the current node itself, and finally the right subtree of the current node recursively.

The postorder traversal is to visit the left subtree of the current node first, then the right subtree of the current node, and finally the current node itself.



Tree Traversal, cont.

The preorder traversal is to visit the current node first, then the left subtree of the current node recursively, and finally the right subtree of the current node recursively.



Tree Traversal, cont.

The breadth-first traversal is to visit the nodes level by level. First visit the root, then all children of the root from left to right, then grandchildren of the root from left to right, and so on.

For example, in the tree in Figure 25.2, the inorder is 45 55 57 59 60 67 100 101 107. The postorder is 45 59 57 55 67 101 107 100 60. The preorder is 60 55 45 57 59 100 67 107 101. The breadth-first traversal is 60 55 100 45 57 67 107 59 101.



The Tree Interface

«interface»
java.lang.Collection<E>



«interface»
Tree<E>

```
+search(e: E): boolean
+insert(e: E): boolean
+delete(e: E): boolean

+inorder(): void
+preorder(): void
+postorder(): void
+getSize(): int
+isEmpty(): boolean
+clear(): void
```

Override the add, isEmpty, remove, containsAll, addAll, removeAll, retainAll, toArray(), and toArray(T[]) methods defined in Collection using default methods.

The Tree interface defines common operations for trees.

Returns true if the specified element is in the tree.

Returns true if the element is added successfully.

Returns true if the element is removed from the tree successfully.

Prints the nodes in inorder traversal.

Prints the nodes in preorder traversal.

Prints the nodes in postorder traversal.

Returns the number of elements in the tree.

Returns true if the tree is empty.

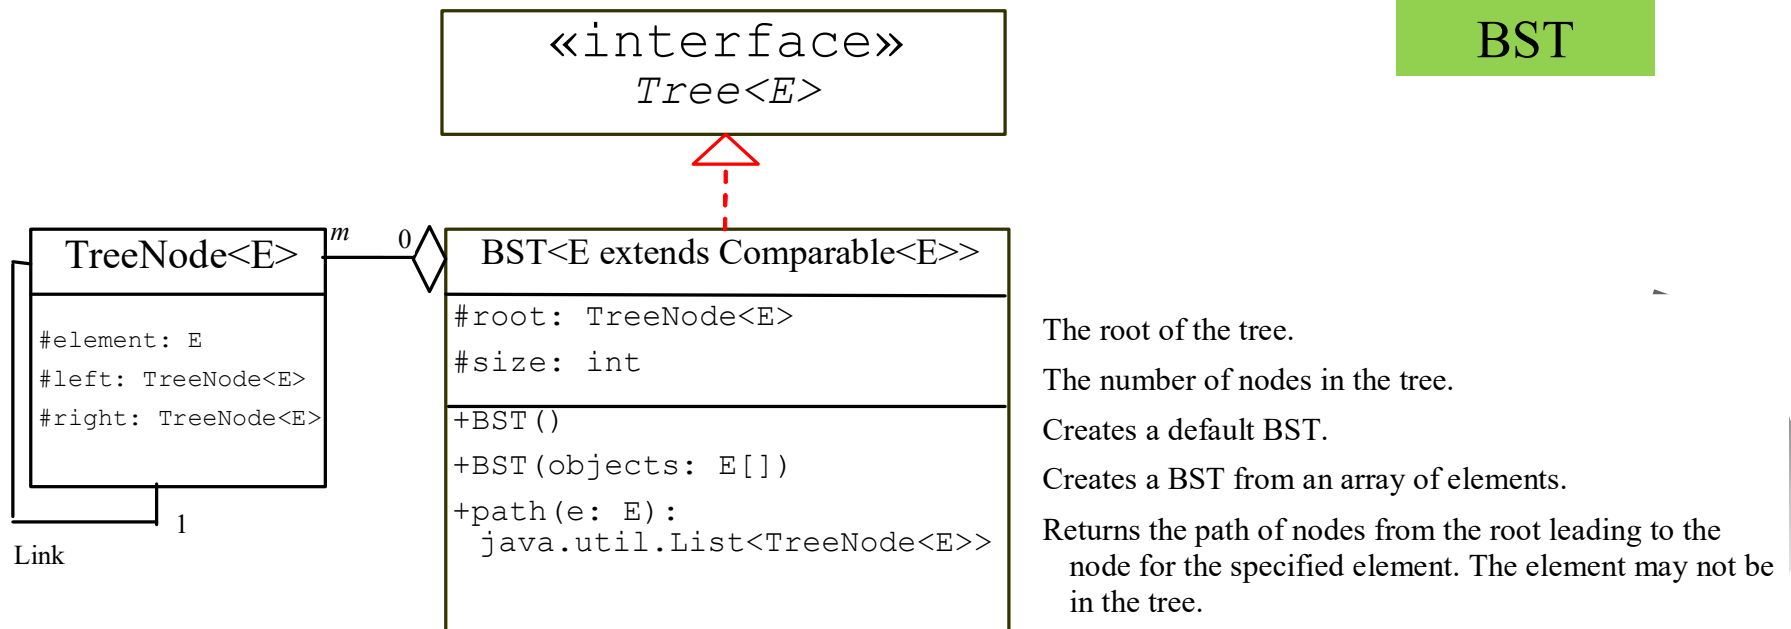
Removes all elements from the tree.

Tree



The BST Class

Let's define the binary tree class, named BST with A concrete BST class can be defined to extend AbstractTree.



Example: Using Binary Trees


Write a program that creates a binary tree using BST. Add strings into the binary tree and traverse the tree in inorder, postorder, and preorder.

TestBST

Run



Tree After Insertions




Inorder: Adam, Daniel
George, Jones, Michael,
Peter, Tom

Postorder: Daniel Adam,
Jones, Peter, Tom,
Michael, George

Preorder: George, Adam,
Daniel, Michael, Jones,
Tom, Peter


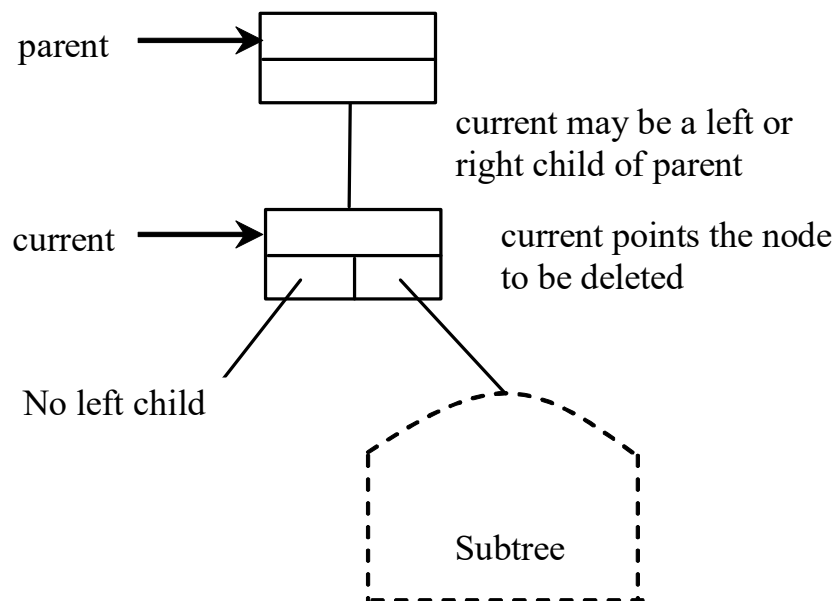
Deleting Elements in a Binary Search Tree

To delete an element from a binary tree, you need to first locate the node that contains the element and also its parent node. Let current point to the node that contains the element in the binary tree and parent point to the parent of the current node. The current node may be a left child or a right child of the parent node. There are two cases to consider:





Deleting Elements in a Binary Search Tree

Case 1: The current node does not have a left child, as shown in this figure (a). Simply connect the parent with the right child of the current node, as shown in this figure (b).




Deleting Elements in a Binary Search Tree

For example, to delete node 10 in Figure 25.9a. Connect the parent of node 10 with the right child of node 10, as shown in Figure 25.9b.



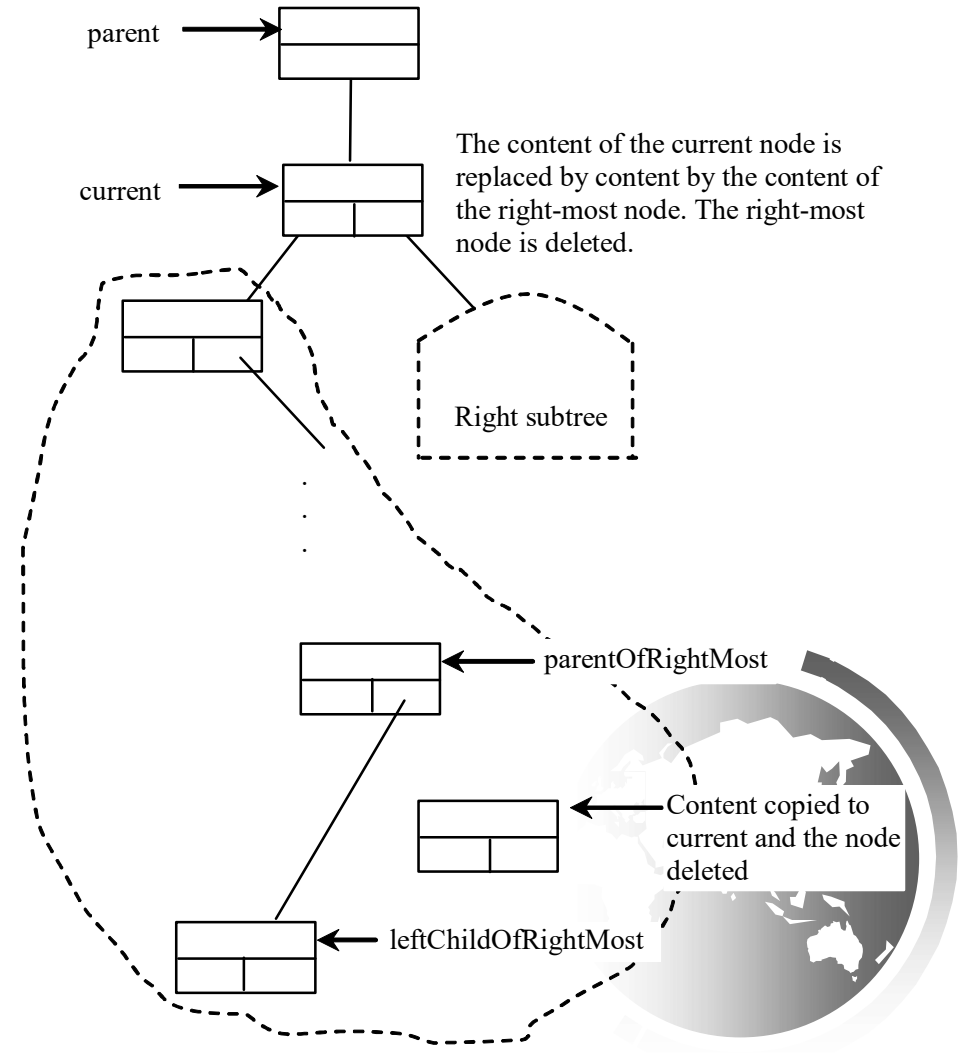
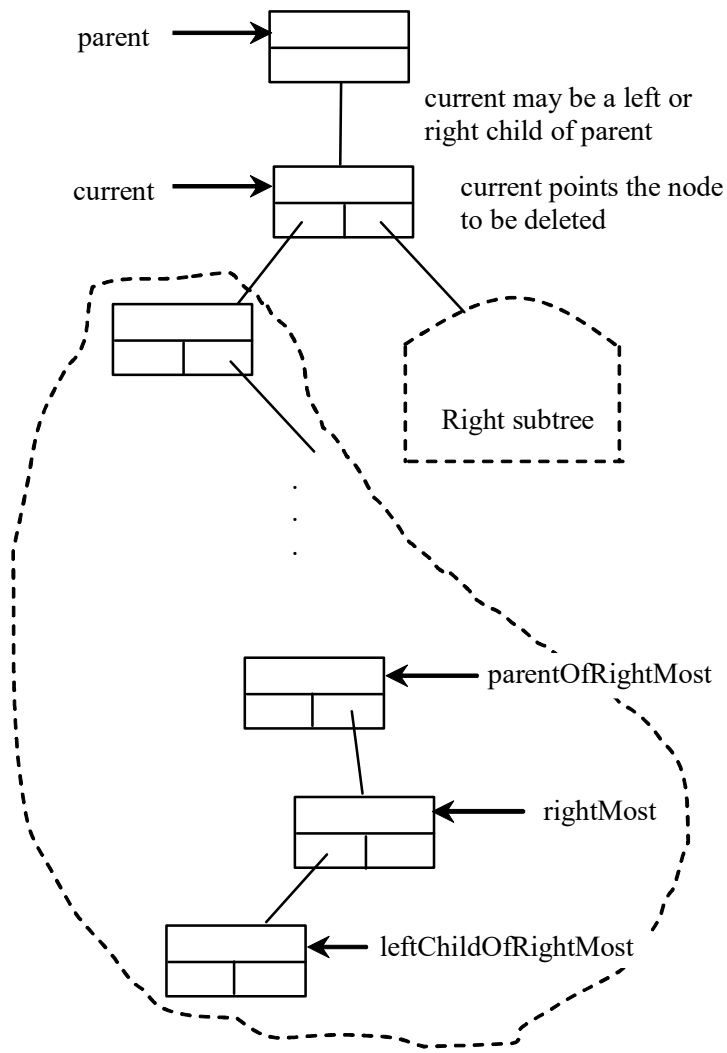
Deleting Elements in a Binary Search Tree

Case 2: The current node has a left child. Let `rightMost` point to the node that contains the largest element in the left subtree of the current node and `parentOfRightMost` point to the parent node of the `rightMost` node, as shown in Figure 25.10a. Note that the `rightMost` node cannot have a right child, but may have a left child. Replace the element value in the current node with the one in the `rightMost` node, connect the `parentOfRightMost` node with the left child of the `rightMost` node, and delete the `rightMost` node, as shown in Figure 25.10b.



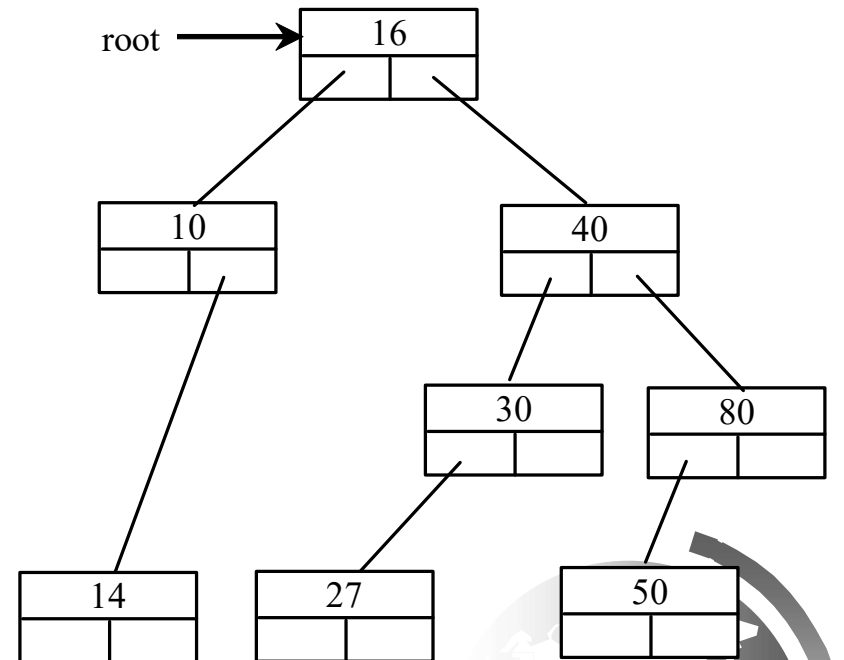

Deleting Elements in a Binary Search Tree

Case 2 diagram


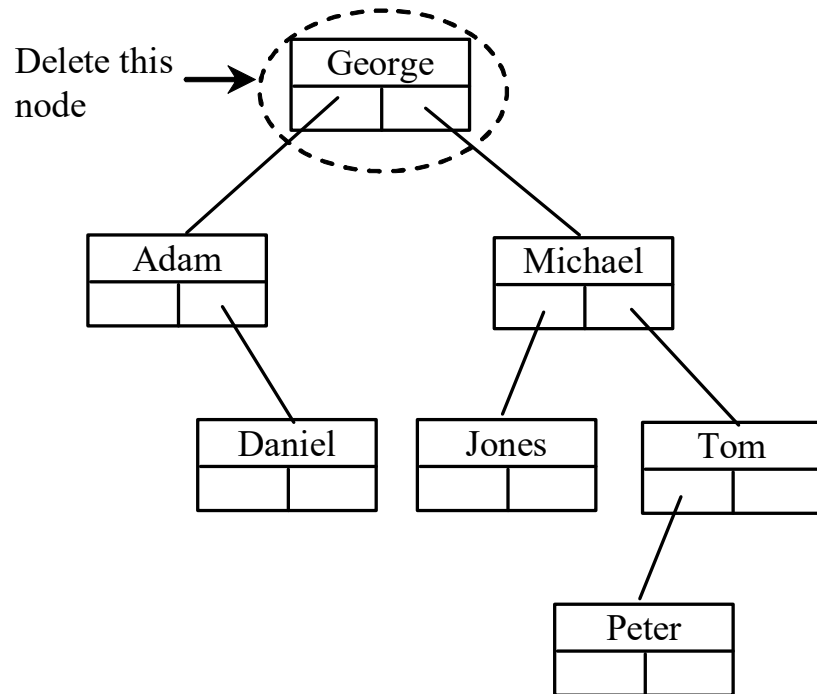


Deleting Elements in a Binary Search Tree



Case 2 example, delete 20





Examples



Examples



Examples



TestBSTDelete

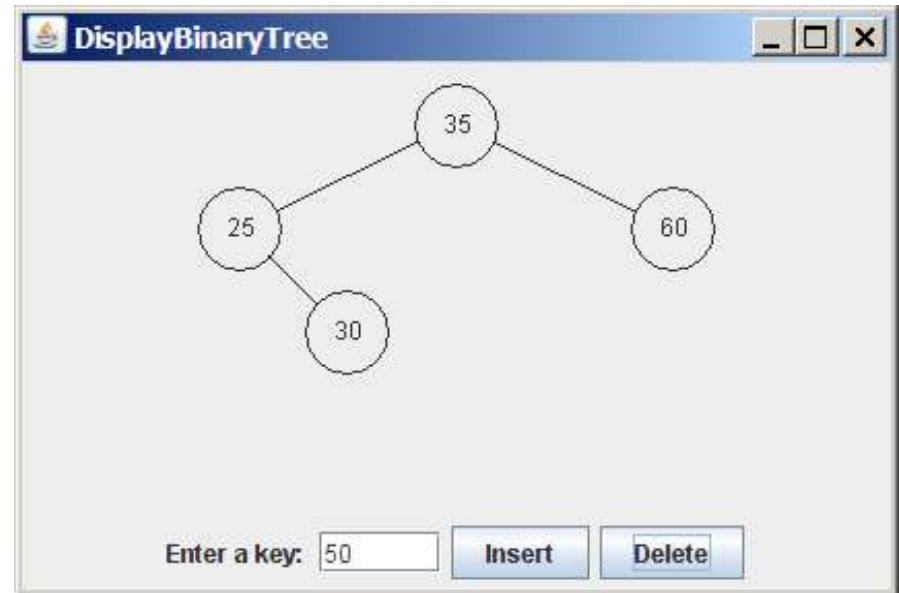
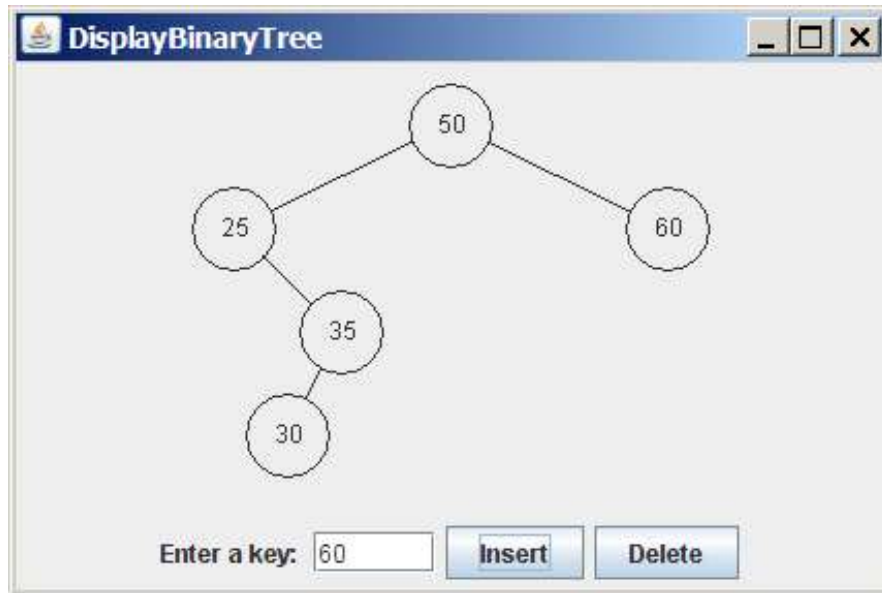
Run

binary tree time complexity

It is obvious that the time complexity for the inorder, preorder, and postorder is $O(n)$, since each node is traversed only once. The time complexity for search, insertion and deletion is the height of the tree. In the worst case, the height of the tree is $O(n)$.



Tree Visualization



BTView

BSTAnimation

Run

Iterators

An *iterator* is an object that provides a uniform way for traversing the elements in a container such as a set, list, binary tree, etc.

```
«interface»  
java.util.Iterator<E>
```

```
+hasNext(): boolean  
+next(): E  
+remove(): void
```

Returns true if the iterator has more elements.

Returns the next element in the iterator.

Removes from the underlying container the last element returned by the iterator (optional operation).


TestBSTWithIterator

Run

Data Compression: Huffman Coding

In ASCII, every character is encoded in 8 bits. Huffman coding compresses data by using fewer bits to encode more frequently occurring characters. The codes for characters are constructed based on the occurrence of characters in the text using a binary tree, called the *Huffman coding tree*.

Mississippi



Character	Code	Frequency
M	000	1
p	001	2
s	01	4
i	1	4

000101011010110010011

21 bits

See How a Binary Tree Works

www.cs.armstrong.edu/liang/animation/HuffmanCodingAnimation.html

The screenshot shows a Mozilla Firefox browser window displaying the "Huffman Coding Animation" web application by Y. Daniel Liang. The application interface includes a text input field containing "Welcome", a "Show Huffman Tree" button, a binary tree diagram, a bit string input field containing "0100110101001010", and a "Decode Text" button. A modal dialog box titled "Encode Text to Bits" is open, displaying the message "Welcome is encoded to 110100111110001010" and an "OK" button. The binary tree diagram shows a root node with value 7, branching into nodes 3 and 4. Node 3 branches into nodes 1 (labeled 'o') and 2 (labeled 'm' and 'l'). Node 4 branches into nodes 2 (labeled 'e') and 2 (labeled 'w' and 'c').

Constructing Huffman Tree

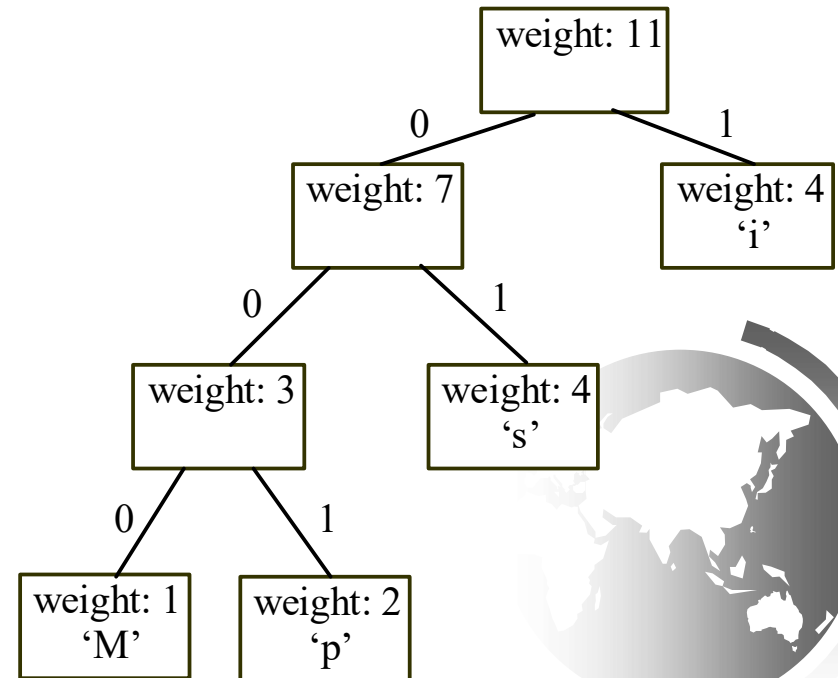



To construct a *Huffman coding tree*, use a greedy algorithm as follows:

- Begin with a forest of trees. Each tree contains a node for a character. The weight of the node is the frequency of the character in the text.
- Repeat this step until there is only one tree:
Choose two trees with the smallest weight and create a new node as their parent. The weight of the new tree is the sum of the weight of the subtrees.



Constructing Huffman Tree

Mississippi



Constructing Huffman Tree

HuffmanCode

Run